

LDCT Denoising With REDCNN & WGAN-GP

Hanbie Ryu

Dept. of A.I.

2024 Medisys Intern



Abstract

- LDCT (Low-dose CT) images require denoising systems to enhance image quality
- Used REDCNN (Residual Encoder-Decoder CNN) for primary model architecture
- Experimented with WGAN-GP (Wasserstein GAN + Gradient Penalty) with simple CNN discriminator for REDCNN generator model training procedure
- Expanded image dataset with variable doses for generalized denoising across different noise levels

REDCNN

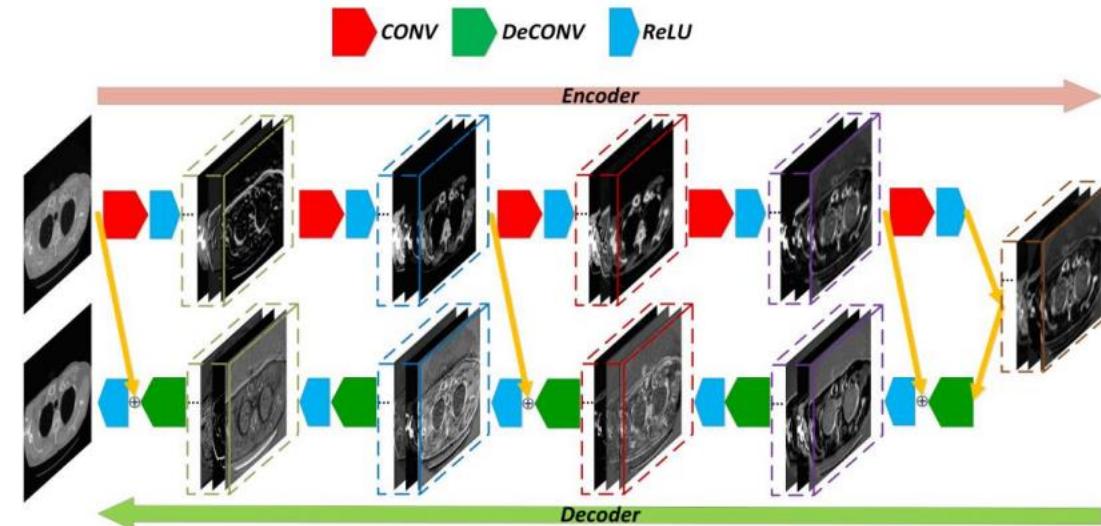


Fig.1. Overall architecture of our proposed RED-CNN network.

“Low-Dose CT with a Residual Encoder-Decoder Convolutional Neural Network (RED-CNN)”
Hu Chen, et al., 2017

Theory

- Residual connections for improved learning process
- Lack of max-pooling layers to minimize down-sampling
- Symmetrical convolutional & deconvolutional layers

Images from Hu Chen, et. al., 2017

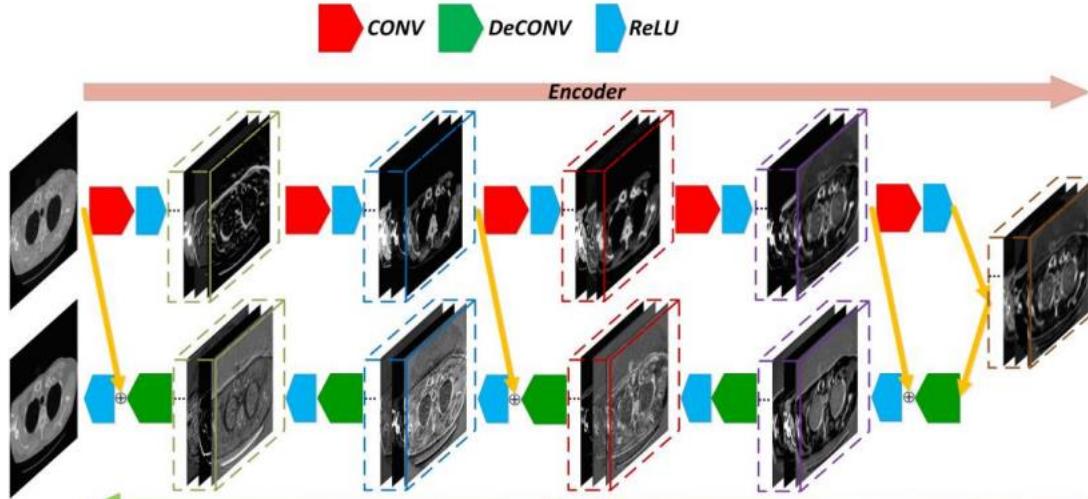


Fig.1. Overall architecture of our proposed RED-CNN network.

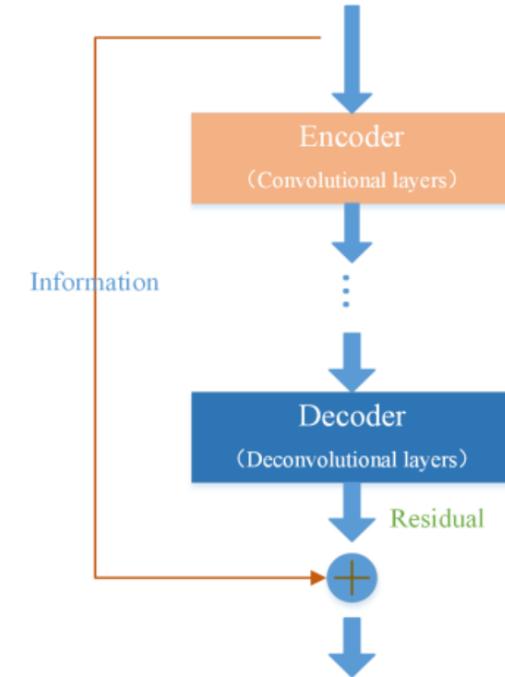


Fig.2. Shortcut in the residual compensation structure.

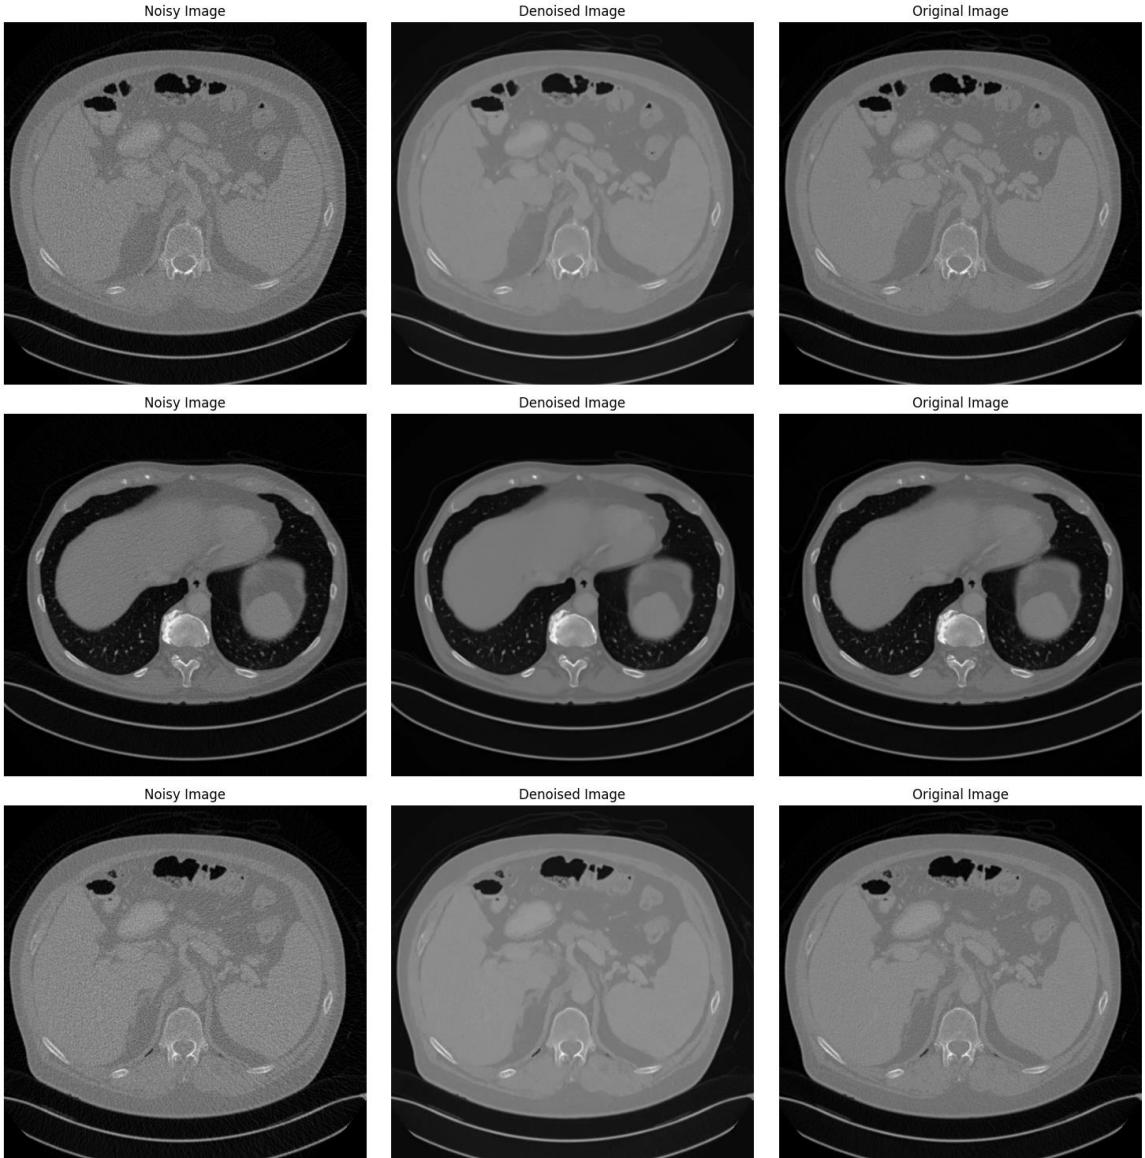
Simple NN training with REDCNN architecture

- 32 epochs*
- MSE loss
- Adam optimizer
- L2 regularization
- LR decay with patience=4, factor=0.5
- Early stopping with patience=6

* Further experiments revealed the network could have used less epochs (16), but early stopping has been implemented here so not a big issue with regularization in this case

REDCNN model analysis

- Rotation of random degrees for augmentation found to be ineffective, unless of 90-degree units
- Followed hyperparameters and model structure from the IEEE 2017 article, except for the patching stride of 16 (small changes seemed to have no particular significance)
- Excellent in terms of pure noise reduction



Limitations

- Built to minimize RMSE loss from target NDCT image
 - => Results in over-smoothing (blurry outputs)
 - => Possibility of blurring out crucial details in CT image
 - => Recognized the need to have a regularising procedure to have the model output images similar to the NDCT images



WGAN-GP

$$L = \underbrace{\mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})] - \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})]}_{\text{Original critic loss}} + \lambda \underbrace{\mathbb{E}_{\hat{\mathbf{x}} \sim \mathbb{P}_{\hat{\mathbf{x}}}} [(\|\nabla_{\hat{\mathbf{x}}} D(\hat{\mathbf{x}})\|_2 - 1)^2]}_{\text{Our gradient penalty}}.$$

[GAN — Wasserstein GAN & WGAN-GP | by Jonathan Hui | Medium](#)

Theory

- GAN (Generative Adversarial Networks)
 - A training process that consists of a generator and a discriminator, where:
 - The generator takes the LDCT image as input (in this application)
 - The generator outputs a fake image
 - The discriminator scores the output image in range $[0, 1] \Rightarrow$ classification
 - The discriminator and generator's loss is calculated from the discriminator's classification output (the generator also takes the regular MSE loss from the NDCT image, weighted between 0 ~ 1)
- ⇒ The resulting images should look more like the NDCT images in theory

Theory

- WGAN-GP (Wasserstein GAN + Gradient penalty)
 - The discriminator outputs a ‘realness’ score with no fixed range, instead of a classification score between [0, 1]
⇒ Improves gradient feedback to the generator
 - Gradient penalty (GP) is applied instead of weight clipping (WGAN-GP) [1]
⇒ Improves optimization while maintaining Lipschitz constraint [2]

$$[1]: \lambda \cdot E_{\hat{x} \sim P_{\hat{x}}}[(\| \nabla_{\hat{x}} f(\hat{x}) \|_2 - 1)^2]$$

$$[2]: |f(x_1) - f(x_2)| \leq K \cdot \|x_1 - x_2\|$$

```
g_loss = disc_weight * -torch.mean(fake_validity) + mse_weight * MSE(gen_images, images)

disc_weight = 0.95
mse_weight = 0.05
```

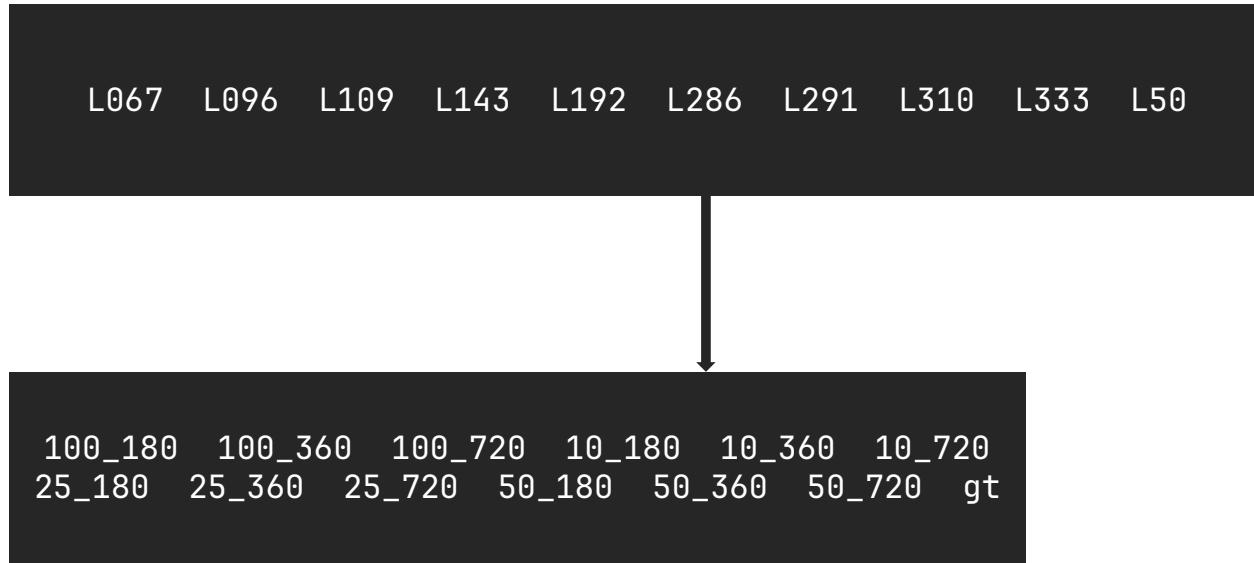
Discriminator model

- Simple CNN architecture with down-sampling convolution layers
- Used LeakyReLU for better gradient flow

```
self.model = nn.Sequential(  
    nn.Conv2d(1, 64, kernel_size=3, stride=1, padding=1),  
    nn.LeakyReLU(0.2),  
    nn.Conv2d(64, 128, kernel_size=5, stride=3, padding=1),  
    nn.LeakyReLU(0.2),  
    nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1),  
    nn.LeakyReLU(0.2),  
    nn.Conv2d(128, 128, kernel_size=5, stride=3, padding=1),  
    nn.LeakyReLU(0.2),  
    nn.Conv2d(128, 1, kernel_size=3, stride=1, padding=0)  
)
```

Variable dose dataset

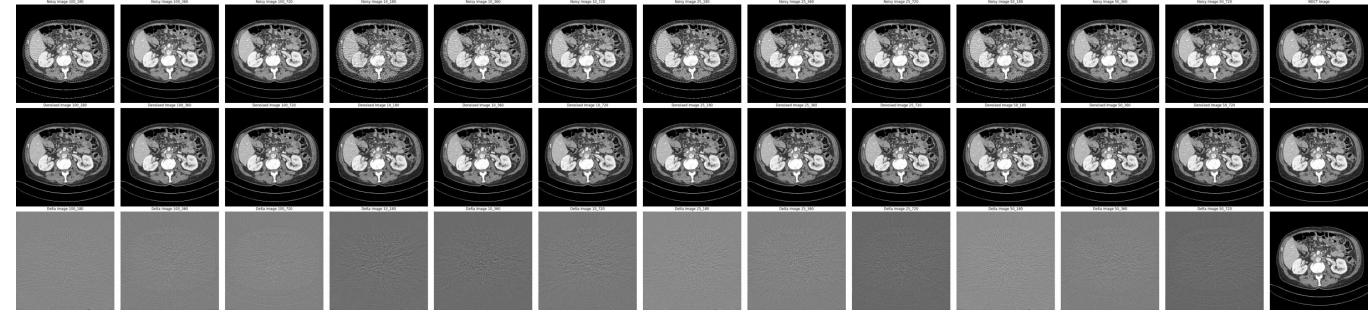
- Trained on dataset of variable noise levels
⇒ Trained for generalized denoising



Dataset source: idk

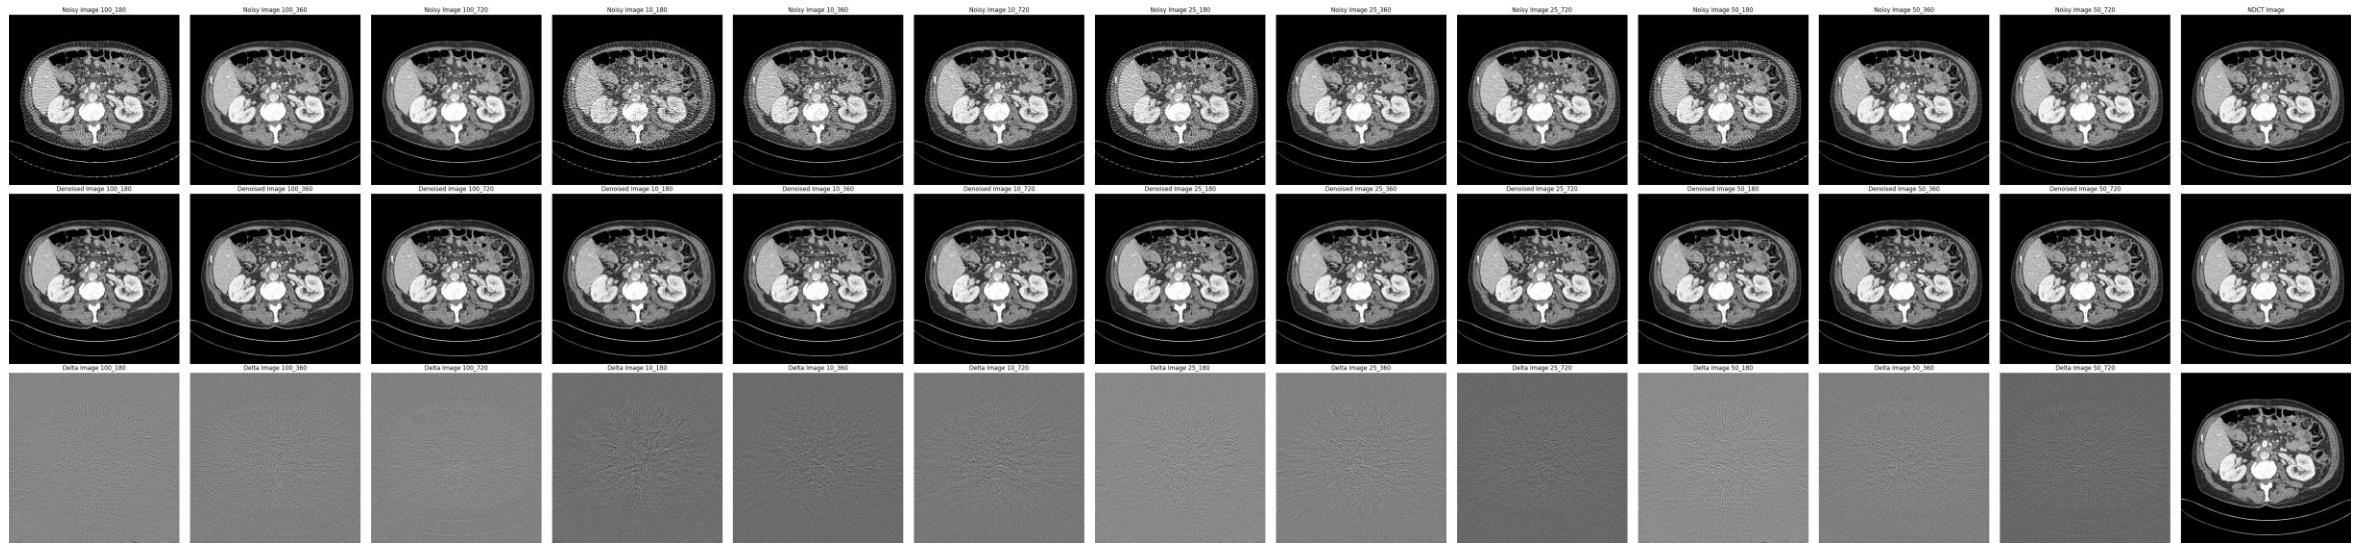
Results

3



Effect of WGAN-GP training process

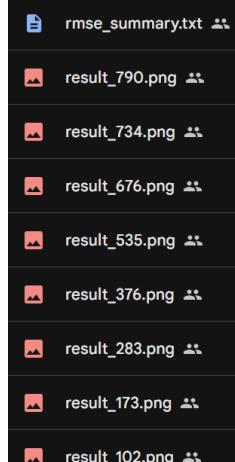
- Images generated by the WGAN-GP models appear to have noise-levels similar to the NDCT target images => retains image details, no over-smoothing
- Pure RMSE scores higher (lower loss) in regular REDCNN models with no GAN training
- Some artifacts ('streaks') remain in very noisy image inputs



Final RMSE scores

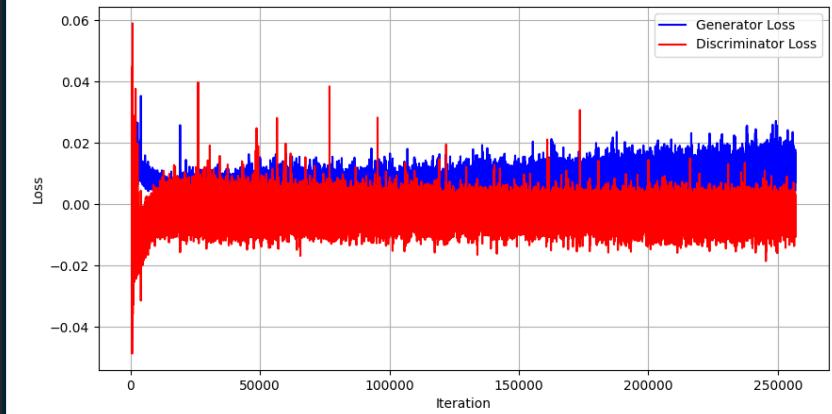
Dose	Average RMSE
100_180	24.04
100_360	20.18
100_720	17.67
10_180	30.61
10_360	25.66
10_720	23.25
25_180	26.77
25_360	22.86
25_720	20.95
50_180	25.18
50_360	21.34
50_720	19.44
mean:	23.1625

Denoised results for 8 randomly selected images
[Drive Link here](#)



```
[Config.yaml]
dataset:
    augment: true
    batch_size: 64
    patch_size: 64
    stride: 32
training:
    betas:
        - 0.9
        - 0.999
    critic_iterations: 5
    l2_regularization: 2e-5
    lambda_gp: 10.0
    learning_rate: 1e-4
    lr_decay_factor: 0.5
    lr_decay_patience: 4
    num_epochs: 16
    patience: 6
    sample_interval: 5
weights:
    disc_weight: 0.05
    mse_weight: 0.95
```

Generator and discriminator loss
(from iteration=500)



Moving forward

- Model variants
 - U-net, U-former etc
- GAN variants
 - ESRGAN
- Restormer?
- To be decided
- 좋은 자원을 제공해주시고, 지도해 주셔서 감사합니다. 앞으로도 최선을 다하겠습니다.